

10 minutes to pandas

This is a short introduction to pandas, geared mainly for new users. You can see more complex recipes in the [Cookbook](#).

Customarily, we import as follows:

```
In [1]: import numpy as np
```

```
In [2]: import pandas as pd
```

Basic data structures in pandas

Pandas provides two types of classes for handling data:

1. **[Series](#)**: a one-dimensional labeled array holding data of any type such as integers, strings, Python objects etc.
2. **[DataFrame](#)**: a two-dimensional data structure that holds data like a two-dimension array or a table with rows and columns.

Object creation

See the [Intro to data structures section](#).

Creating a [Series](#) by passing a list of values, letting pandas create a default [RangeIndex](#).

```
In [3]: s = pd.Series([1, 3, 5, np.nan, 6, 8])
```

```
In [4]: s
```

```
Out[4]:
```

```
0    1.0
```

```
1    3.0
```

```
2    5.0
```

```
3    NaN
```

```
4    6.0
```

```
5    8.0
```

```
dtype: float64
```

Creating a [DataFrame](#) by passing a NumPy array with a datetime index using [date_range\(\)](#) and labeled columns:

```
In [5]: dates = pd.date_range("20130101", periods=6)
```

```
In [6]: dates
```

```
Out[6]:
```

```
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
               '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
```

```
In [7]: df = pd.DataFrame(np.random.randn(6, 4), index=dates,
                          columns=list("ABCD"))
```

```
In [8]: df
```

```
Out[8]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Creating a [DataFrame](#) by passing a dictionary of objects where the keys are the column labels and the values are the column values.

```
In [9]: df2 = pd.DataFrame(
...:     {
...:         "A": 1.0,
...:         "B": pd.Timestamp("20130102"),
...:         "C": pd.Series(1, index=list(range(4)), dtype="float32"),
...:         "D": np.array([3] * 4, dtype="int32"),
...:         "E": pd.Categorical(["test", "train", "test", "train"]),
...:         "F": "foo",
...:     }
...: )
...:
```

```
In [10]: df2
```

```
Out[10]:
```

	A	B	C	D	E	F
0	1.0	2013-01-02	1.0	3	test	foo
1	1.0	2013-01-02	1.0	3	train	foo
2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

The columns of the resulting [DataFrame](#) have different [dtypes](#):

```
In [11]: df2.dtypes
```

```
Out[11]:
```

A	float64
B	datetime64[s]
C	float32
D	int32
E	category

```
F          object
dtype: object
```

If you're using IPython, tab completion for column names (as well as public attributes) is automatically enabled. Here's a subset of the attributes that will be completed:

```
In [12]: df2.<TAB> # noqa: E225, E999
df2.A          df2.bool
df2.abs         df2.boxplot
df2.add         df2.C
df2.add_prefix df2.clip
df2.add_suffix  df2.columns
df2.align       df2.copy
df2.all         df2.count
df2.any         df2.combine
df2.append      df2.D
df2.apply       df2.describe
df2.applymap    df2.diff
df2.B          df2.duplicated
```

As you can see, the columns A, B, C, and D are automatically tab completed. E and F are there as well; the rest of the attributes have been truncated for brevity.

Viewing data

See the [Essentially basics functionality section](#).

Use [DataFrame.head\(\)](#) and [DataFrame.tail\(\)](#) to view the top and bottom rows of the frame respectively:

```
In [13]: df.head()
Out[13]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

```
In [14]: df.tail(3)
Out[14]:
```

	A	B	C	D
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

Display the [DataFrame.index](#) or [DataFrame.columns](#):

```
In [15]: df.index
```

```
Out[15]:
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',
              '2013-01-05', '2013-01-06'],
              dtype='datetime64[ns]', freq='D')
```

```
In [16]: df.columns
Out[16]: Index(['A', 'B', 'C', 'D'], dtype='object')
```

Return a NumPy representation of the underlying data with **DataFrame.to_numpy()** without the index or column labels:

```
In [17]: df.to_numpy()
Out[17]:
array([[ 0.4691, -0.2829, -1.5091, -1.1356],
       [ 1.2121, -0.1732,  0.1192, -1.0442],
       [-0.8618, -2.1046, -0.4949,  1.0718],
       [ 0.7216, -0.7068, -1.0396,  0.2719],
       [-0.425 ,  0.567 ,  0.2762, -1.0874],
       [-0.6737,  0.1136, -1.4784,  0.525 ]])
```

Note

NumPy arrays have one dtype for the entire array while pandas DataFrames have one dtype per column. When you call **DataFrame.to_numpy()**, pandas will find the NumPy dtype that can hold *all* of the dtypes in the DataFrame. If the common data type is object, **DataFrame.to_numpy()** will require copying data.

```
In [18]: df2.dtypes
Out[18]:
A          float64
B    datetime64[s]
C          float32
D           int32
E          category
F           object
dtype: object
```

```
In [19]: df2.to_numpy()
Out[19]:
array([[1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'test', 'foo'],
       [1.0, Timestamp('2013-01-02 00:00:00'), 1.0, 3, 'train', 'foo']],
      dtype=object)
```

[describe\(\)](#) shows a quick statistic summary of your data:

```
In [20]: df.describe()
Out[20]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	0.073711	-0.431125	-0.687758	-0.233103
std	0.843157	0.922818	0.779887	0.973118
min	-0.861849	-2.104569	-1.509059	-1.135632
25%	-0.611510	-0.600794	-1.368714	-1.076610
50%	0.022070	-0.228039	-0.767252	-0.386188
75%	0.658444	0.041933	-0.034326	0.461706
max	1.212112	0.567020	0.276232	1.071804

Transposing your data:

```
In [21]: df.T
```

```
Out[21]:
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04	2013-01-05	2013-01-06
A	0.469112	1.212112	-0.861849	0.721555	-0.424972	-0.673690
B	-0.282863	-0.173215	-2.104569	-0.706771	0.567020	0.113648
C	-1.509059	0.119209	-0.494929	-1.039575	0.276232	-1.478427
D	-1.135632	-1.044236	1.071804	0.271860	-1.087401	0.524988

[`DataFrame.sort\_index\(\)`](#) sorts by an axis:

```
In [22]: df.sort_index(axis=1, ascending=False)
```

```
Out[22]:
```

	D	C	B	A
2013-01-01	-1.135632	-1.509059	-0.282863	0.469112
2013-01-02	-1.044236	0.119209	-0.173215	1.212112
2013-01-03	1.071804	-0.494929	-2.104569	-0.861849
2013-01-04	0.271860	-1.039575	-0.706771	0.721555
2013-01-05	-1.087401	0.276232	0.567020	-0.424972
2013-01-06	0.524988	-1.478427	0.113648	-0.673690

[`DataFrame.sort\_values\(\)`](#) sorts by values:

```
In [23]: df.sort_values(by="B")
```

```
Out[23]:
```

	A	B	C	D
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-06	-0.673690	0.113648	-1.478427	0.524988
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

Selection

Note

While standard Python / NumPy expressions for selecting and setting are intuitive and come in handy for interactive work, for production code, we recommend the optimized pandas data access methods, [DataFrame.at\(\)](#), [DataFrame.iat\(\)](#), [DataFrame.loc\(\)](#) and [DataFrame.iloc\(\)](#).

See the indexing documentation [Indexing and Selecting Data](#) and [MultiIndex / Advanced Indexing](#).

Getitem ([])

For a [DataFrame](#), passing a single label selects a columns and yields a [Series](#) equivalent to `df.A`:

```
In [24]: df["A"]
Out[24]:
2013-01-01    0.469112
2013-01-02    1.212112
2013-01-03   -0.861849
2013-01-04    0.721555
2013-01-05   -0.424972
2013-01-06   -0.673690
Freq: D, Name: A, dtype: float64
```

For a [DataFrame](#), passing a slice : selects matching rows:

```
In [25]: df[0:3]
Out[25]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804

```
In [26]: df["20130102":"20130104"]
Out[26]:
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selection by label

See more in [Selection by Label](#) using [DataFrame.loc\(\)](#) or [DataFrame.at\(\)](#).

Selecting a row matching a label:

```
In [27]: df.loc[dates[0]]
Out[27]:
```

	A
A	0.469112
B	-0.282863
C	-1.509059

```
D    -1.135632
Name: 2013-01-01 00:00:00, dtype: float64
```

Selecting all rows (:) with a select column labels:

```
In [28]: df.loc[:, ["A", "B"]]
```

```
Out[28]:
```

	A	B
2013-01-01	0.469112	-0.282863
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771
2013-01-05	-0.424972	0.567020
2013-01-06	-0.673690	0.113648

For label slicing, both endpoints are *included*:

```
In [29]: df.loc["20130102":"20130104", ["A", "B"]]
```

```
Out[29]:
```

	A	B
2013-01-02	1.212112	-0.173215
2013-01-03	-0.861849	-2.104569
2013-01-04	0.721555	-0.706771

Selecting a single row and column label returns a scalar:

```
In [30]: df.loc[dates[0], "A"]
```

```
Out[30]: 0.4691122999071863
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [31]: df.at[dates[0], "A"]
```

```
Out[31]: 0.4691122999071863
```

Selection by position

See more in [Selection by Position](#) using [DataFrame.iloc\(\)](#) or [DataFrame.iat\(\)](#).

Select via the position of the passed integers:

```
In [32]: df.iloc[3]
```

```
Out[32]:
```

	A	B	C	D
	0.721555	-0.706771	-1.039575	0.271860

```
Name: 2013-01-04 00:00:00, dtype: float64
```

Integer slices acts similar to NumPy/Python:

```
In [33]: df.iloc[3:5, 0:2]
Out[33]:
```

	A	B
2013-01-04	0.721555	-0.706771
2013-01-05	-0.424972	0.567020

Lists of integer position locations:

```
In [34]: df.iloc[[1, 2, 4], [0, 2]]
Out[34]:
```

	A	C
2013-01-02	1.212112	0.119209
2013-01-03	-0.861849	-0.494929
2013-01-05	-0.424972	0.276232

For slicing rows explicitly:

```
In [35]: df.iloc[1:3, :]
Out[35]:
```

	A	B	C	D
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804

For slicing columns explicitly:

```
In [36]: df.iloc[:, 1:3]
Out[36]:
```

	B	C
2013-01-01	-0.282863	-1.509059
2013-01-02	-0.173215	0.119209
2013-01-03	-2.104569	-0.494929
2013-01-04	-0.706771	-1.039575
2013-01-05	0.567020	0.276232
2013-01-06	0.113648	-1.478427

For getting a value explicitly:

```
In [37]: df.iloc[1, 1]
Out[37]: -0.17321464905330858
```

For getting fast access to a scalar (equivalent to the prior method):

```
In [38]: df.iat[1, 1]
Out[38]: -0.17321464905330858
```

Boolean indexing

Select rows where df.A is greater than 0.

```
In [39]: df[df["A"] > 0]
```


Out[39]:

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-04	0.721555	-0.706771	-1.039575	0.271860

Selecting values from a [DataFrame](#) where a boolean condition is met:

```
In [40]: df[df > 0]
```

Out[40]:

	A	B	C	D
2013-01-01	0.469112	NaN	NaN	NaN
2013-01-02	1.212112	NaN	0.119209	NaN
2013-01-03	NaN	NaN	NaN	1.071804
2013-01-04	0.721555	NaN	NaN	0.271860
2013-01-05	NaN	0.567020	0.276232	NaN
2013-01-06	NaN	0.113648	NaN	0.524988

Using [isin\(\)](#) method for filtering:

```
In [41]: df2 = df.copy()
```

```
In [42]: df2["E"] = ["one", "one", "two", "three", "four", "three"]
```

```
In [43]: df2
```

Out[43]:

	A	B	C	D	E
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632	one
2013-01-02	1.212112	-0.173215	0.119209	-1.044236	one
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804	two
2013-01-04	0.721555	-0.706771	-1.039575	0.271860	three
2013-01-05	-0.424972	0.567020	0.276232	-1.087401	four
2013-01-06	-0.673690	0.113648	-1.478427	0.524988	three

```
In [44]: df2[df2["E"].isin(["two", "four"])]
```

Out[44]:

	A	B	C	D	E
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804	two
2013-01-05	-0.424972	0.567020	0.276232	-1.087401	four

Setting

Setting a new column automatically aligns the data by the indexes:

```
In [45]: s1 = pd.Series([1, 2, 3, 4, 5, 6],  
index=pd.date_range("20130102", periods=6))
```

```
In [46]: s1
```

Out[46]:

```
2013-01-02    1
2013-01-03    2
2013-01-04    3
2013-01-05    4
2013-01-06    5
2013-01-07    6
Freq: D, dtype: int64
```

```
In [47]: df["F"] = s1
```

Setting values by label:

```
In [48]: df.at[dates[0], "A"] = 0
```

Setting values by position:

```
In [49]: df.iat[0, 1] = 0
```

Setting by assigning with a NumPy array:

```
In [50]: df.loc[:, "D"] = np.array([5] * len(df))
```

The result of the prior setting operations:

```
In [51]: df
```

```
Out[51]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-1.509059	5.0	NaN
2013-01-02	1.212112	-0.173215	0.119209	5.0	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5.0	2.0
2013-01-04	0.721555	-0.706771	-1.039575	5.0	3.0
2013-01-05	-0.424972	0.567020	0.276232	5.0	4.0
2013-01-06	-0.673690	0.113648	-1.478427	5.0	5.0

A where operation with setting:

```
In [52]: df2 = df.copy()
```

```
In [53]: df2[df2 > 0] = -df2
```

```
In [54]: df2
```

```
Out[54]:
```

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-1.509059	-5.0	NaN
2013-01-02	-1.212112	-0.173215	-0.119209	-5.0	-1.0
2013-01-03	-0.861849	-2.104569	-0.494929	-5.0	-2.0
2013-01-04	-0.721555	-0.706771	-1.039575	-5.0	-3.0
2013-01-05	-0.424972	-0.567020	-0.276232	-5.0	-4.0
2013-01-06	-0.673690	-0.113648	-1.478427	-5.0	-5.0

Missing data

For NumPy data types, `np.nan` represents missing data. It is by default not included in computations. See the [Missing Data section](#).

Reindexing allows you to change/add/delete the index on a specified axis. This returns a copy of the data:

```
In [55]: df1 = df.reindex(index=dates[0:4], columns=list(df.columns) + ["E"])
```

```
In [56]: df1.loc[dates[0] : dates[1], "E"] = 1
```

```
In [57]: df1
```

```
Out[57]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-1.509059	5.0	NaN	1.0
2013-01-02	1.212112	-0.173215	0.119209	5.0	1.0	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5.0	2.0	NaN
2013-01-04	0.721555	-0.706771	-1.039575	5.0	3.0	NaN

[`DataFrame.dropna\(\)`](#) drops any rows that have missing data:

```
In [58]: df1.dropna(how="any")
```

```
Out[58]:
```

	A	B	C	D	F	E
2013-01-02	1.212112	-0.173215	0.119209	5.0	1.0	1.0

[`DataFrame.fillna\(\)`](#) fills missing data:

```
In [59]: df1.fillna(value=5)
```

```
Out[59]:
```

	A	B	C	D	F	E
2013-01-01	0.000000	0.000000	-1.509059	5.0	5.0	1.0
2013-01-02	1.212112	-0.173215	0.119209	5.0	1.0	1.0
2013-01-03	-0.861849	-2.104569	-0.494929	5.0	2.0	5.0
2013-01-04	0.721555	-0.706771	-1.039575	5.0	3.0	5.0

[`isna\(\)`](#) gets the boolean mask where values are nan:

```
In [60]: pd.isna(df1)
```

```
Out[60]:
```

	A	B	C	D	F	E
2013-01-01	False	False	False	False	True	False
2013-01-02	False	False	False	False	False	False
2013-01-03	False	False	False	False	False	True
2013-01-04	False	False	False	False	False	True

Operations

See the [Basic section on Binary Ops](#).

Stats

Operations in general *exclude* missing data.

Calculate the mean value for each column:

```
In [61]: df.mean()
Out[61]:
A    -0.004474
B    -0.383981
C    -0.687758
D     5.000000
F     3.000000
dtype: float64
```

Calculate the mean value for each row:

```
In [62]: df.mean(axis=1)
Out[62]:
2013-01-01    0.872735
2013-01-02    1.431621
2013-01-03    0.707731
2013-01-04    1.395042
2013-01-05    1.883656
2013-01-06    1.592306
Freq: D, dtype: float64
```

Operating with another [Series](#) or [DataFrame](#) with a different index or column will align the result with the union of the index or column labels. In addition, pandas automatically broadcasts along the specified dimension and will fill unaligned labels with `np.nan`.

```
In [63]: s = pd.Series([1, 3, 5, np.nan, 6, 8], index=dates).shift(2)
```

```
In [64]: s
Out[64]:
2013-01-01    NaN
2013-01-02    NaN
2013-01-03     1.0
2013-01-04     3.0
2013-01-05     5.0
2013-01-06    NaN
Freq: D, dtype: float64
```

```
In [65]: df.sub(s, axis="index")
```

Out[65]:

	A	B	C	D	F
2013-01-01	NaN	NaN	NaN	NaN	NaN
2013-01-02	NaN	NaN	NaN	NaN	NaN
2013-01-03	-1.861849	-3.104569	-1.494929	4.0	1.0
2013-01-04	-2.278445	-3.706771	-4.039575	2.0	0.0
2013-01-05	-5.424972	-4.432980	-4.723768	0.0	-1.0
2013-01-06	NaN	NaN	NaN	NaN	NaN

User defined functions

[`DataFrame.agg\(\)`](#) and [`DataFrame.transform\(\)`](#) applies a user defined function that reduces or broadcasts its result respectively.

```
In [66]: df.agg(lambda x: np.mean(x) * 5.6)
```

Out[66]:

```
A    -0.025054
B    -2.150294
C    -3.851445
D     28.000000
F     16.800000
dtype: float64
```

```
In [67]: df.transform(lambda x: x * 101.2)
```

Out[67]:

	A	B	C	D	F
2013-01-01	0.000000	0.000000	-152.716721	506.0	NaN
2013-01-02	122.665737	-17.529322	12.063922	506.0	101.2
2013-01-03	-87.219115	-212.982405	-50.086843	506.0	202.4
2013-01-04	73.021382	-71.525239	-105.204988	506.0	303.6
2013-01-05	-43.007200	57.382459	27.954680	506.0	404.8
2013-01-06	-68.177398	11.501219	-149.616767	506.0	506.0

Value Counts

See more at [Histogramming and Discretization](#).

```
In [68]: s = pd.Series(np.random.randint(0, 7, size=10))
```

```
In [69]: s
```

Out[69]:

```
0    4
1    2
2    1
3    2
4    6
5    4
6    4
```

```

7      6
8      4
9      4
dtype: int64

In [70]: s.value_counts()
Out[70]:
4      5
2      2
6      2
1      1
Name: count, dtype: int64

```

String Methods

[Series](#) is equipped with a set of string processing methods in the `str` attribute that make it easy to operate on each element of the array, as in the code snippet below. See more at [Vectorized String Methods](#).

```

In [71]: s = pd.Series(["A", "B", "C", "Aaba", "Baca", np.nan, "CABA",
                        "dog", "cat"])

In [72]: s.str.lower()
Out[72]:
0      a
1      b
2      c
3    aaba
4    baca
5     NaN
6    caba
7     dog
8     cat
dtype: object

```

Merge

Concat

pandas provides various facilities for easily combining together [Series](#) and [DataFrame](#) objects with various kinds of set logic for the indexes and relational algebra functionality in the case of join / merge-type operations.

See the [Merging section](#).

Concatenating pandas objects together row-wise with [concat\(\)](#):

```
In [73]: df = pd.DataFrame(np.random.randn(10, 4))
```

```
In [74]: df
```

```
Out[74]:
```

	0	1	2	3
0	-0.548702	1.467327	-1.015962	-0.483075
1	1.637550	-1.217659	-0.291519	-1.745505
2	-0.263952	0.991460	-0.919069	0.266046
3	-0.709661	1.669052	1.037882	-1.705775
4	-0.919854	-0.042379	1.247642	-0.009920
5	0.290213	0.495767	0.362949	1.548106
6	-1.131345	-0.089329	0.337863	-0.945867
7	-0.932132	1.956030	0.017587	-0.016692
8	-0.575247	0.254161	-1.143704	0.215897
9	1.193555	-0.077118	-0.408530	-0.862495

```
# break it into pieces
```

```
In [75]: pieces = [df[:3], df[3:7], df[7:]]
```

```
In [76]: pd.concat(pieces)
```

```
Out[76]:
```

	0	1	2	3
0	-0.548702	1.467327	-1.015962	-0.483075
1	1.637550	-1.217659	-0.291519	-1.745505
2	-0.263952	0.991460	-0.919069	0.266046
3	-0.709661	1.669052	1.037882	-1.705775
4	-0.919854	-0.042379	1.247642	-0.009920
5	0.290213	0.495767	0.362949	1.548106
6	-1.131345	-0.089329	0.337863	-0.945867
7	-0.932132	1.956030	0.017587	-0.016692
8	-0.575247	0.254161	-1.143704	0.215897
9	1.193555	-0.077118	-0.408530	-0.862495

Note

Adding a column to a [DataFrame](#) is relatively fast. However, adding a row requires a copy, and may be expensive. We recommend passing a pre-built list of records to the [DataFrame](#) constructor instead of building a [DataFrame](#) by iteratively appending records to it.

Join

[merge\(\)](#) enables SQL style join types along specific columns. See the [Database style joining](#) section.

```
In [77]: left = pd.DataFrame({"key": ["foo", "foo"], "lval": [1, 2]})
```

```
In [78]: right = pd.DataFrame({"key": ["foo", "foo"], "rval": [4, 5]})
```

```

In [79]: left
Out[79]:
   key  lval
0  foo     1
1  foo     2

In [80]: right
Out[80]:
   key  rval
0  foo     4
1  foo     5

In [81]: pd.merge(left, right, on="key")
Out[81]:
   key  lval  rval
0  foo     1     4
1  foo     1     5
2  foo     2     4
3  foo     2     5

```

[merge\(\)](#) on unique keys:

```

In [82]: left = pd.DataFrame({"key": ["foo", "bar"], "lval": [1, 2]})

In [83]: right = pd.DataFrame({"key": ["foo", "bar"], "rval": [4, 5]})

In [84]: left
Out[84]:
   key  lval
0  foo     1
1  bar     2

In [85]: right
Out[85]:
   key  rval
0  foo     4
1  bar     5

In [86]: pd.merge(left, right, on="key")
Out[86]:
   key  lval  rval
0  foo     1     4
1  bar     2     5

```

Grouping

By “group by” we are referring to a process involving one or more of the following steps:

- **Splitting** the data into groups based on some criteria
- **Applying** a function to each group independently
- **Combining** the results into a data structure

See the [Grouping section](#).

```
In [87]: df = pd.DataFrame(
.....:     {
.....:         "A": ["foo", "bar", "foo", "bar", "foo", "bar", "foo",
"foo"],
.....:         "B": ["one", "one", "two", "three", "two", "two", "one",
"three"],
.....:         "C": np.random.randn(8),
.....:         "D": np.random.randn(8),
.....:     }
.....: )
.....:
```

```
In [88]: df
```

```
Out[88]:
```

	A	B	C	D
0	foo	one	1.346061	-1.577585
1	bar	one	1.511763	0.396823
2	foo	two	1.627081	-0.105381
3	bar	three	-0.990582	-0.532532
4	foo	two	-0.441652	1.453749
5	bar	two	1.211526	1.208843
6	foo	one	0.268520	-0.080952
7	foo	three	0.024580	-0.264610

Grouping by a column label, selecting column labels, and then applying the [DataFrameGroupBy.sum\(\)](#) function to the resulting groups:

```
In [89]: df.groupby("A")[["C", "D"]].sum()
Out[89]:
```

		C	D
A			
bar		1.732707	1.073134
foo		2.824590	-0.574779

Grouping by multiple columns label forms [MultiIndex](#).

```
In [90]: df.groupby(["A", "B"]).sum()
Out[90]:
```

			C	D
A	B			
bar	one		1.511763	0.396823
	three		-0.990582	-0.532532
	two		1.211526	1.208843
foo	one		1.614581	-1.658537

```
three  0.024580 -0.264610
two    1.185429  1.348368
```

Reshaping

See the sections on [Hierarchical Indexing](#) and [Reshaping](#).

Stack

```
In [91]: arrays = [
.....:     ["bar", "bar", "baz", "baz", "foo", "foo", "qux", "qux"],
.....:     ["one", "two", "one", "two", "one", "two", "one", "two"],
.....: ]
.....:
```

```
In [92]: index = pd.MultiIndex.from_arrays(arrays, names=["first",
"second"])
```

```
In [93]: df = pd.DataFrame(np.random.randn(8, 2), index=index,
columns=["A", "B"])
```

```
In [94]: df2 = df[:4]
```

```
In [95]: df2
```

```
Out[95]:
```

		A	B
first	second		
bar	one	-0.727965	-0.589346
	two	0.339969	-0.693205
baz	one	-0.339355	0.593616
	two	0.884345	1.591431

The [stack\(\)](#) method “compresses” a level in the DataFrame’s columns:

```
In [96]: stacked = df2.stack(future_stack=True)
```

```
In [97]: stacked
```

```
Out[97]:
```

first	second		
bar	one	A	-0.727965
		B	-0.589346
	two	A	0.339969
		B	-0.693205
baz	one	A	-0.339355
		B	0.593616
	two	A	0.884345
		B	1.591431

dtype: float64

With a “stacked” DataFrame or Series (having a [MultiIndex](#) as the index), the inverse operation of `stack()` is `unstack()`, which by default unstacks the **last level**:

```
In [98]: stacked.unstack()
```

```
Out[98]:
```

		A	B
first	second		
bar	one	-0.727965	-0.589346
	two	0.339969	-0.693205
baz	one	-0.339355	0.593616
	two	0.884345	1.591431

```
In [99]: stacked.unstack(1)
```

```
Out[99]:
```

		one	two
first			
bar	A	-0.727965	0.339969
	B	-0.589346	-0.693205
baz	A	-0.339355	0.884345
	B	0.593616	1.591431

```
In [100]: stacked.unstack(0)
```

```
Out[100]:
```

		bar	baz
first			
second			
one	A	-0.727965	-0.339355
	B	-0.589346	0.593616
two	A	0.339969	0.884345
	B	-0.693205	1.591431

Pivot tables

See the section on [Pivot Tables](#).

```
In [101]: df = pd.DataFrame(  
.....:     {  
.....:         "A": ["one", "one", "two", "three"] * 3,  
.....:         "B": ["A", "B", "C"] * 4,  
.....:         "C": ["foo", "foo", "foo", "bar", "bar", "bar"] * 2,  
.....:         "D": np.random.randn(12),  
.....:         "E": np.random.randn(12),  
.....:     }  
.....: )  
.....:
```

```
In [102]: df
```

```
Out[102]:
```

	A	B	C	D	E
--	---	---	---	---	---

```

0      one  A  foo -1.202872  0.047609
1      one  B  foo -1.814470 -0.136473
2      two  C  foo  1.018601 -0.561757
3  three  A  bar -0.595447 -1.623033
4      one  B  bar  1.395433  0.029399
5      one  C  bar -0.392670 -0.542108
6      two  A  foo  0.007207  0.282696
7  three  B  foo  1.928123 -0.087302
8      one  C  foo -0.055224 -1.575170
9      one  A  bar  2.395985  1.771208
10     two  B  bar  1.552825  0.816482
11    three  C  bar  0.166599  1.100230

```

[`pivot\_table\(\)`](#) pivots a [DataFrame](#) specifying the values, index and columns

```

In [103]: pd.pivot_table(df, values="D", index=["A", "B"], columns=["C"])
Out[103]:

```

		bar	foo
one	A	2.395985	-1.202872
	B	1.395433	-1.814470
	C	-0.392670	-0.055224
three	A	-0.595447	NaN
	B	NaN	1.928123
	C	0.166599	NaN
two	A	NaN	0.007207
	B	1.552825	NaN
	C	NaN	1.018601

Time series

pandas has simple, powerful, and efficient functionality for performing resampling operations during frequency conversion (e.g., converting secondly data into 5-minutely data). This is extremely common in, but not limited to, financial applications. See the [Time Series section](#).

```

In [104]: rng = pd.date_range("1/1/2012", periods=100, freq="s")

In [105]: ts = pd.Series(np.random.randint(0, 500, len(rng)), index=rng)

In [106]: ts.resample("5Min").sum()
Out[106]:
2012-01-01    24182
Freq: 5min, dtype: int64

```

[`Series.tz\_localize\(\)`](#) localizes a time series to a time zone:

```

In [107]: rng = pd.date_range("3/6/2012 00:00", periods=5, freq="D")

```

```
In [108]: ts = pd.Series(np.random.randn(len(rng)), rng)
```

```
In [109]: ts
```

```
Out[109]:
2012-03-06    1.857704
2012-03-07   -1.193545
2012-03-08    0.677510
2012-03-09   -0.153931
2012-03-10    0.520091
Freq: D, dtype: float64
```

```
In [110]: ts_utc = ts.tz_localize("UTC")
```

```
In [111]: ts_utc
```

```
Out[111]:
2012-03-06 00:00:00+00:00    1.857704
2012-03-07 00:00:00+00:00   -1.193545
2012-03-08 00:00:00+00:00    0.677510
2012-03-09 00:00:00+00:00   -0.153931
2012-03-10 00:00:00+00:00    0.520091
Freq: D, dtype: float64
```

[Series.tz_convert\(\)](#) converts a timezones aware time series to another time zone:

```
In [112]: ts_utc.tz_convert("US/Eastern")
```

```
Out[112]:
2012-03-05 19:00:00-05:00    1.857704
2012-03-06 19:00:00-05:00   -1.193545
2012-03-07 19:00:00-05:00    0.677510
2012-03-08 19:00:00-05:00   -0.153931
2012-03-09 19:00:00-05:00    0.520091
Freq: D, dtype: float64
```

Adding a non-fixed duration ([BusinessDay](#)) to a time series:

```
In [113]: rng
```

```
Out[113]:
DatetimeIndex(['2012-03-06', '2012-03-07', '2012-03-08', '2012-03-09',
               '2012-03-10'],
              dtype='datetime64[ns]', freq='D')
```

```
In [114]: rng + pd.offsets.BusinessDay(5)
```

```
Out[114]:
DatetimeIndex(['2012-03-13', '2012-03-14', '2012-03-15', '2012-03-16',
               '2012-03-16'],
              dtype='datetime64[ns]', freq=None)
```

Categoricals

pandas can include categorical data in a [DataFrame](#). For full docs, see the [categorical introduction](#) and the [API documentation](#).

```
In [115]: df = pd.DataFrame(
.....:     {"id": [1, 2, 3, 4, 5, 6], "raw_grade": ["a", "b", "b", "a",
"a", "e"]}
.....: )
.....:
```

Converting the raw grades to a categorical data type:

```
In [116]: df["grade"] = df["raw_grade"].astype("category")
```

```
In [117]: df["grade"]
Out[117]:
0      a
1      b
2      b
3      a
4      a
5      e
Name: grade, dtype: category
Categories (3, object): ['a', 'b', 'e']
```

Rename the categories to more meaningful names:

```
In [118]: new_categories = ["very good", "good", "very bad"]
```

```
In [119]: df["grade"] = df["grade"].cat.rename_categories(new_categories)
```

Reorder the categories and simultaneously add the missing categories (methods under [Series.cat\(\)](#) return a new [Series](#) by default):

```
In [120]: df["grade"] = df["grade"].cat.set_categories(
.....:     ["very bad", "bad", "medium", "good", "very good"]
.....: )
.....:

In [121]: df["grade"]
Out[121]:
0      very good
1           good
2           good
3      very good
4      very good
5      very bad
Name: grade, dtype: category
Categories (5, object): ['very bad', 'bad', 'medium', 'good', 'very good']
```

Sorting is per order in the categories, not lexical order:

```
In [122]: df.sort_values(by="grade")
```

```
Out[122]:
```

	id	raw_grade	grade
5	6	e	very bad
1	2	b	good
2	3	b	good
0	1	a	very good
3	4	a	very good
4	5	a	very good

Grouping by a categorical column with `observed=False` also shows empty categories:

```
In [123]: df.groupby("grade", observed=False).size()
```

```
Out[123]:
```

grade	
very bad	1
bad	0
medium	0
good	2
very good	3

dtype: int64

Plotting

See the [Plotting](#) docs.

We use the standard convention for referencing the matplotlib API:

```
In [124]: import matplotlib.pyplot as plt
```


```
In [125]: plt.close("all")
```

The `plt.close` method is used to [close](#) a figure window:

```
In [126]: ts = pd.Series(np.random.randn(1000),  
index=pd.date_range("1/1/2000", periods=1000))
```

```
In [127]: ts = ts.cumsum()
```

```
In [128]: ts.plot();
```



Note

When using Jupyter, the plot will appear using [plot\(\)](#). Otherwise use [matplotlib.pyplot.show](#) to show it or [matplotlib.pyplot.savefig](#) to write it to a file.

[plot\(\)](#) plots all columns:


```
In [129]: df = pd.DataFrame(  
.....:     np.random.randn(1000, 4), index=ts.index, columns=["A", "B",  
"C", "D"]  
.....: )  
.....:
```

```
In [130]: df = df.cumsum()
```

```
In [131]: plt.figure();
```

```
In [132]: df.plot();
```

```
In [133]: plt.legend(loc='best');
```

Importing and exporting data

See the [IO Tools](#) section.

CSV

Writing to a csv file: using [DataFrame.to_csv\(\)](#)

```
In [134]: df = pd.DataFrame(np.random.randint(0, 5, (10, 5)))
```

```
In [135]: df.to_csv("foo.csv")
```

Reading from a csv file: using [read_csv\(\)](#)

```
In [136]: pd.read_csv("foo.csv")
```

```
Out[136]:
```

```
   Unnamed: 0  0  1  2  3  4
0            0  4  3  1  1  2
1            1  1  0  2  3  2
2            2  1  4  2  1  2
```

```

3          3  0  4  0  2  2
4          4  4  2  2  3  4
5          5  4  0  4  3  1
6          6  2  1  2  0  3
7          7  4  0  4  4  4
8          8  4  4  1  0  1
9          9  0  4  3  0  3

```

Parquet

Writing to a Parquet file:

```
In [137]: df.to_parquet("foo.parquet")
```

Reading from a Parquet file Store using [read_parquet\(\)](#):

```
In [138]: pd.read_parquet("foo.parquet")
```

```
Out[138]:
```

```

   0  1  2  3  4
0  4  3  1  1  2
1  1  0  2  3  2
2  1  4  2  1  2
3  0  4  0  2  2
4  4  2  2  3  4
5  4  0  4  3  1
6  2  1  2  0  3
7  4  0  4  4  4
8  4  4  1  0  1
9  0  4  3  0  3

```

Excel

Reading and writing to [Excel](#).

Writing to an excel file using [DataFrame.to_excel\(\)](#):

```
In [139]: df.to_excel("foo.xlsx", sheet_name="Sheet1")
```

Reading from an excel file using [read_excel\(\)](#):

```
In [140]: pd.read_excel("foo.xlsx", "Sheet1", index_col=None,
na_values=["NA"])
```

```
Out[140]:
```

```

   Unnamed: 0  0  1  2  3  4
0            0  4  3  1  1  2
1            1  1  0  2  3  2
2            2  1  4  2  1  2
3            3  0  4  0  2  2

```

4	4	4	2	2	3	4
5	5	4	0	4	3	1
6	6	2	1	2	0	3
7	7	4	0	4	4	4
8	8	4	4	1	0	1
9	9	0	4	3	0	3

Gotchas

If you are attempting to perform a boolean operation on a [Series](#) or [DataFrame](#) you might see an exception like:

```
In [141]: if pd.Series([False, True, False]):
.....:     print("I was true")
.....:

-----
-
ValueError                                Traceback (most recent call
last)
<ipython-input-141-b27eb9c1dfc0> in ?()
----> 1 if pd.Series([False, True, False]):
      2     print("I was true")

~/work/pandas/pandas/pandas/core/generic.py in ?(self)
    1574     @final
    1575     def __nonzero__(self) -> NoReturn:
-> 1576         raise ValueError(
    1577             f"The truth value of a {type(self).__name__} is
ambiguous. "
    1578             "Use a.empty, a.bool(), a.item(), a.any() or a.all()."
    1579         )

ValueError: The truth value of a Series is ambiguous. Use a.empty,
a.bool(), a.item(), a.any() or a.all().
```

See [Comparisons](#) and [Gotchas](#) for an explanation and what to do.